

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: data =pd.read_csv("F:/winequality-red.csv")
data
```

Out[2]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	5
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8	5
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8	5
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8	6
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	5
...
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5	5
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11.2	6
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	11.0	6
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10.2	5
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11.0	6

1599 rows × 12 columns

```
In [3]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density                1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol                1599 non-null   float64
11  quality                1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

```
In [4]: data.describe()
```

Out[4]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	s
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	159
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.467792	0.996747	3.311113	
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.895324	0.001887	0.154386	
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000	0.990070	2.740000	
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000	0.995600	3.210000	
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000	0.996750	3.310000	
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000	0.997835	3.400000	
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000	1.003690	4.010000	

```
In [5]: data.isna().sum()
```

```
Out[5]: fixed acidity      0
volatile acidity      0
citric acid           0
residual sugar        0
chlorides             0
free sulfur dioxide    0
total sulfur dioxide   0
density               0
pH                   0
sulphates             0
alcohol               0
quality               0
dtype: int64
```







```
In [6]: data.rename(columns={
    "fixed acidity": "fixed_acidity",
    "volatile acidity": "volatile_acidity",
    "citric acid": "citric_acid",
    "residual sugar": "residual_sugar",
    "chlorides": "chlorides",
    "free sulfur dioxide": "free_sulfur_dioxide",
    "total sulfur dioxide": "total_sulfur_dioxide"
},inplace=True)
```













```
In [7]: data.columns
```





```
Out[7]: Index(['fixed_acidity', 'volatile_acidity', 'citric_acid', 'residual_sugar',
              'chlorides', 'free_sulfur_dioxide', 'total_sulfur_dioxide', 'density',
              'pH', 'sulphates', 'alcohol', 'quality'],
              dtype='object')
```

```
In [8]: columns=list(data.columns)
```


```
In [9]: fig,ax =plt.subplots(11,2,figsize=(15,45))
plt.subplots_adjust(hspace=0.5)
for i in range(11):
    sns.boxplot(x=columns[i],data=data,ax=ax[i,0])
    sns.scatterplot(x=columns[i],y="quality",data=data,hue="quality",ax=ax[i,1])
```








```
In [10]: corr=data.corr()
sns.heatmap(corr,annot=True,linewidth=0.5,mask=np.triu(corr))
plt.show()
```




```
In [11]: sns.pairplot(data,hue="quality",corner=True)
```


```
Out[11]: <seaborn.axisgrid.PairGrid at 0x1e0449879d0>
```




```
In [12]: correlation_matrix = data.corr(method='spearman')
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.show()
```




```
In [13]: correlation_matrix = data.corr(method='kendall')
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.show()
```



```
In [14]: correlation_matrix = data.corr(method='pearson')
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.show()
```

```
In [15]: sns.scatterplot(x='citric_acid', y='quality', data=data)
plt.show()
```



```
In [16]: # Pearson Correlation
pearson_corr = data['fixed_acidity'].corr(data['citric_acid'])
print(f'Pearson Correlation: {pearson_corr}')


# Spearman Correlation
spearman_corr = data['fixed_acidity'].corr(data['citric_acid'], method='spearman')
print(f'Spearman Correlation: {spearman_corr}')

# Kendall Correlation
kendall_corr = data['fixed_acidity'].corr(data['citric_acid'], method='kendall')
print(f'Kendall Correlation: {kendall_corr}')
```

```
Pearson Correlation: 0.6717034347641059
Spearman Correlation: 0.6617084166678848
Kendall Correlation: 0.484271229113174
```

```
In [17]: sns.scatterplot(data=columns)
```

```
Out[17]: <Axes: >
```



```
In [18]: data.quality.unique()
```

```
Out[18]: array([5, 6, 7, 4, 8, 3], dtype=int64)
```

```
In [19]: data=data.replace({"quality":{8 : 'Good',
                                         7 : 'Good',
                                         6 : 'Middle',
                                         5 : 'Middle',
                                         4 : 'Bad',
                                         3 : 'Bad',}}})
```

```
In [20]: data
```

```
Out[20]:
```

	fixed_acidity	volatile_acidity	citric_acid	residual_sugar	chlorides	free_sulfur_dioxide	total_sulfur_dioxide	density	pH	quality
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	3
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	3
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	4
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	5
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	3
...
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	6
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	6
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	7
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	7
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	8

1599 rows × 12 columns

```
In [21]: from sklearn.preprocessing import MinMaxScaler
```

```
In [22]: X_temp=data.drop(columns="quality")
          y=data.quality
```

```
In [23]: scaler=MinMaxScaler(feature_range=(0,1)).fit_transform(X_temp)
          X=pd.DataFrame(scaler,columns=X_temp.columns)
```

```
In [24]: X.describe()
```


Out[24]:	fixed_acidity	volatile_acidity	citric_acid	residual_sugar	chlorides	free_sulfur_dioxide	total_sulfur_dioxide	density
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
mean	0.329171	0.279329	0.270976	0.112247	0.125988	0.209506	0.142996	0.490211
std	0.154079	0.122644	0.194801	0.096570	0.078573	0.147326	0.116238	0.138577
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.221239	0.184932	0.090000	0.068493	0.096828	0.084507	0.056537	0.406021
50%	0.292035	0.273973	0.260000	0.089041	0.111853	0.183099	0.113074	0.490451
75%	0.407080	0.356164	0.420000	0.116438	0.130217	0.281690	0.197880	0.570111
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

```
In [25]: from sklearn.model_selection import train_test_split,GridSearchCV,cross_val_score,KFold
from sklearn.ensemble import RandomForestClassifier
from sklearn import metrics
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
```

```
In [ ]:
```

```
In [26]: def confusion_plot(y_test,y_prediction):
cm=metrics.confusion_matrices(y_test,y_prediction)
ax=plt.subplot()
ax = sns.heatmap(cm, annot=True, fmt='', cmap="Purples")
ax.set_xlabel('Prediced labels', fontsize=18)
ax.set_ylabel('True labels', fontsize=18)
ax.set_title('Confusion Matrix', fontsize=25)
ax.xaxis.set_ticklabels(['Bad', 'Good', 'Middle'])
ax.yaxis.set_ticklabels(['Bad', 'Good', 'Middle'])
plt.show()
```

```
In [27]: def clfr_plot(y_test,y_pred):
cr=pd.DataFrame(metrics.Classification_report(y_test,y_pred_rf,digits=3,output_dict=True)).T
cr.drop(columns="support",inplace=True)
sns.heatmap(cr,annot=True,cmap="Purples",linecolor="white",linewidths="0.5").xaxis.tick_top()
```

```
In [28]: def clf_plot(y_pred):
cm=metrics.confussion_matrices(y_test,y_pred)
cr=pd.DataFrame(metrics.classification_report(y_test,y_pred_rf,digits=3,output_dict=True)).T
cr.drop(columns="suport",inplace=True)
fig,ax=plt.subplot(1,2,figsize=(15,5))
ax[0] = sns.heatmap(cm, annot=True, fmt='', cmap="Purples", ax=ax[0])
ax[0].set_xlabel('Prediced labels', fontsize=18)
ax[0].set_ylabel('True labels', fontsize=18)
ax[0].set_title('Confusion Matrix', fontsize=25)
ax[0].xaxis.set_ticklabels(['Bad', 'Good', 'Middle'])
ax[0].yaxis.set_ticklabels(['Bad', 'Good', 'Middle'])

# Right AX : Classification Report
ax[1] = sns.heatmap(cr, cmap='Purples', annot=True, linecolor='white', linewidths=0.5, ax=ax[1])
ax[1].xaxis.tick_top()
ax[1].set_title('Classification Report', fontsize=25)
plt.show()
```

```
In [29]: data.quality.value_counts()
```

```
Out[29]: quality
Middle    1319
Good      217
Bad        63
Name: count, dtype: int64
```

```
In [30]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25,random_state=0)
```

```
In [31]: from sklearn.ensemble import RandomForestClassifier
parameters = {
    'n_estimators': [50, 150, 500],
    'criterion': ['gini', 'entropy', 'log_loss'],
    'max_features': ['sqrt', 'log2']
}

# Initialize the classifier
```

```

rf = RandomForestClassifier(n_jobs=-1)

# Perform grid search
rf_cv = GridSearchCV(estimator=rf, cv=20, param_grid=parameters)
rf_cv.fit(X_train, y_train)

# Print the results
print('Tuned hyper parameters:', rf_cv.best_params_)
print('Accuracy:', rf_cv.best_score_)

```

Tuned hyper parameters: {'criterion': 'gini', 'max_features': 'sqrt', 'n_estimators': 150}
Accuracy: 0.8657203389830508

```

In [32]: X_temp=data.drop(columns="quality")
y=data.quality

```

```

In [33]: scaler=MinMaxScaler(feature_range=(0,1)).fit_transform(X_temp)
X=pd.DataFrame(scaler,columns=X_temp.columns)


```

```

In [34]: X.boxplot()

```

Out[34]: <Axes: >

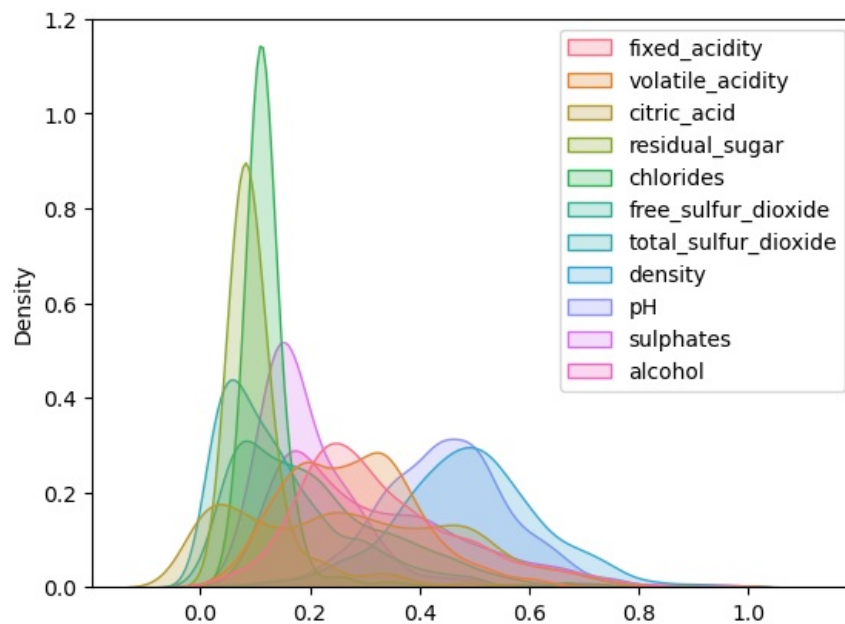


```

In [35]: # KDE plot
sns.kdeplot(data=X, shade=True)

```

Out[35]: <Axes: ylabel='Density'>



```

In [36]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix

```

```

from sklearn.datasets import load_iris

# Load the Iris dataset

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Initialize the RandomForestClassifier
clf = RandomForestClassifier(n_estimators=100, random_state=42)

# Train the classifier
clf.fit(X_train, y_train)

# Make predictions on the test data
y_pred = clf.predict(X_test)

# Evaluate the model
print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

```

```

Classification Report:
              precision    recall  f1-score   support

      Bad           1.00        0.06      0.11         18
      Good           0.60        0.51      0.55         67
      Middle         0.88        0.94      0.91        395

 accuracy              0.85         480
 macro avg              0.83         480
weighted avg              0.85         480

```

```

Confusion Matrix:
[[ 1  0 17]
 [ 0 34 33]
 [ 0 23 372]]

```

In [37]: `X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)`

```

# Initialize the RandomForestClassifier
lr = LogisticRegression( random_state=42)

# Train the classifier
lr.fit(X_train, y_train)

# Make predictions on the test data
y_pred = lr.predict(X_test)

# Evaluate the model
print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

```

```

Classification Report:
              precision    recall  f1-score   support

      Bad           0.00        0.00      0.00         18
      Good           0.59        0.19      0.29         67
      Middle         0.84        0.98      0.91        395

 accuracy              0.83         480
 macro avg              0.48         480
weighted avg              0.78         480

```

```

Confusion Matrix:
[[ 0  0 18]
 [ 0 13 54]
 [ 0  9 386]]

```

In [38]: `# Initialize the Support Vector Classifier`
`svc = SVC(kernel='linear', random_state=42)`

```

# Train the classifier
svc.fit(X_train, y_train)

# Make predictions on the test data
y_pred = svc.predict(X_test)

```

```
# Evaluate the model
print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

```
Classification Report:
              precision    recall  f1-score   support

     Bad         0.00         0.00         0.00         18
     Good         0.00         0.00         0.00         67
    Middle         0.82         1.00         0.90        395

 accuracy         0.27         0.33         0.30        480
 macro avg         0.27         0.33         0.30        480
weighted avg         0.68         0.82         0.74        480

Confusion Matrix:
[[ 0  0 18]
 [ 0  0 67]
 [ 0  0 395]]
```

```
In [39]: from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.metrics import classification_report, confusion_matrix

# Introduce imbalance by removing some instances of one class
X = X[y != 2]
y = y[y != 2]

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Standardize features by removing the mean and scaling to unit variance
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Initialize the Support Vector Classifier with class weight adjustment
svc = SVC(kernel='linear', class_weight='balanced', random_state=42)

# Train the classifier
svc.fit(X_train, y_train)

# Make predictions on the test data
y_pred = svc.predict(X_test)

# Evaluate the model
print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

```
Classification Report:
              precision    recall  f1-score   support

     Bad         0.14         0.89         0.24         18
     Good         0.40         0.84         0.54         67
    Middle         0.96         0.55         0.70        395

 accuracy         0.50         0.76         0.60        480
 macro avg         0.50         0.76         0.50        480
weighted avg         0.85         0.60         0.66        480

Confusion Matrix:
[[ 16  0  2]
 [  4 56  7]
 [ 93 84 218]]
```

```
In [40]: X = X[y != 2]
y = y[y != 2]
# Split data into training and test sets
train_X, test_X, train_y, test_y = train_test_split(X, y, random_state=42, test_size=0.3)

# Scale the data
scaler = StandardScaler()
X_train = scaler.fit_transform(train_X)
X_test = scaler.transform(test_X)

# Train the SVM classifier
```

```

svc = SVC(kernel="linear", class_weight="balanced", random_state=42)
svc.fit(X_train, train_y)

# Predict quality based on user input
try:
    user_input = [float(x) for x in input("Enter values separated by spaces (e.g., '1.5 2.5'): ").split()]
    user_input_scaled = scaler.transform([user_input]) # Scale the input using the same scaler
    prediction = svc.predict(user_input_scaled)
    print("Predicted quality is:", prediction[0])
except Exception as e:
    print("Error:", e)

```

Enter values separated by spaces (e.g., '1.5 2.5'): 7.4 0.700 0.00 1.9 0.076 11.0 34.0 0.99780
 3.51 0.56 9.4
 Predicted quality is: Middle

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js